




Distributed Convex Optimization Framework based on Bulk Synchronous Parallel (BSP) model

Presenter: Behroz Sikander

Supervisor: Prof. Dr. Hans-Arno Jacobsen

Advisor: Dipl.-Ing. Jose Adan Rivera Acevedo


Date: 4th March, 2016



Trend towards **electric vehicles**




Trend towards **electric vehicles**




Electric vehicles takes **more** power than a house





Trend towards **electric vehicles**




Electric vehicles takes **more** power than a house




9 Mio EVs -> Grid overloaded! **Increase Infrastructure.**




Keep infrastructure. **Controlled EV charging**




Keep infrastructure. **Controlled EV charging**




Proposed solution -> **EVADMM**



Keep infrastructure. **Controlled EV charging**



Proposed solution -> **EVADMM**



Thesis Goal: **Distributed implementation & Framework**



Total time to process N EVs using
 M machines ?

Total machines required to process
 N EVs in T time ?

Agenda

- ▷ Background
- ▷ Algorithm
- ▷ Deployment
- ▷ Results
- ▷ Framework

Background

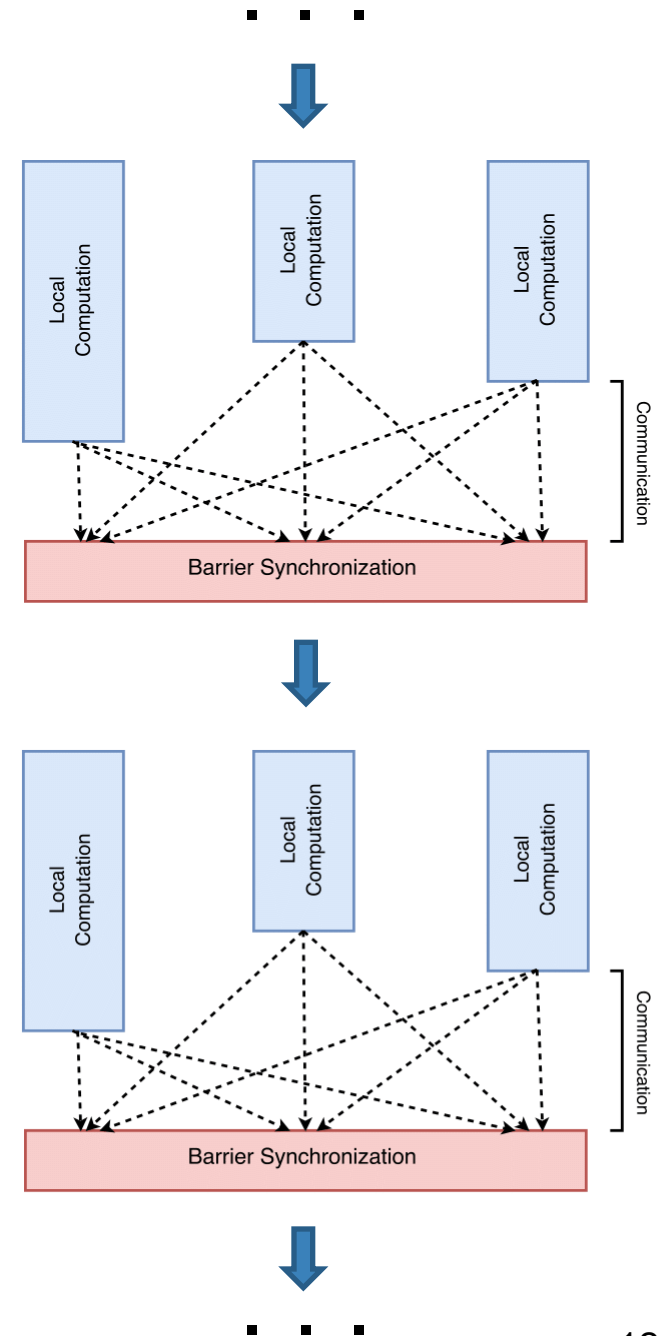
BSP

- ▷ Bulk Synchronous Parallel
- ▷ Developed by Leslie Valiant in 1980s
- ▷ *Model for designing parallel algorithm*
- ▷ Strong theoretical background
- ▷ BSP computer has
 - p processors
 - each with local memory
 - point-point communication

Supersteps

▷ Computation divided in *supersteps*

- concurrent local computation (w)
- global communication (h)
- barrier synchronization (l)



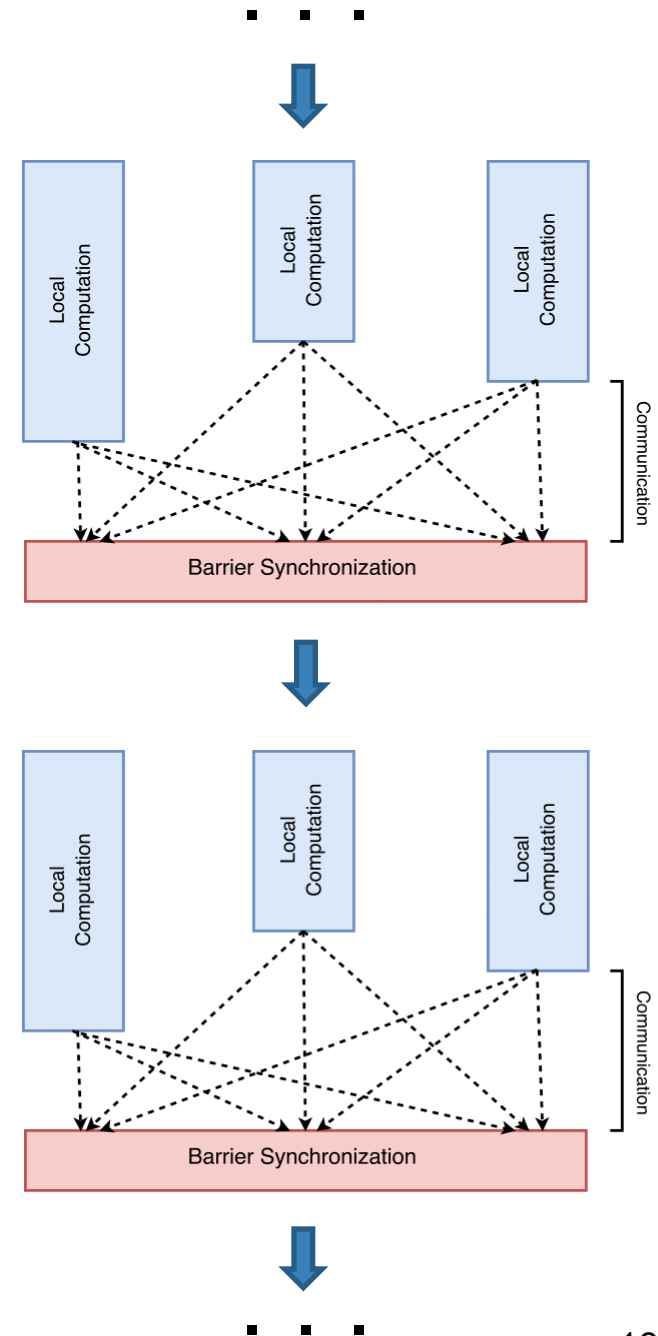
Supersteps

▷ Computation divided in *supersteps*

- concurrent local computation (w)
- global communication (h)
- barrier synchronization (l)

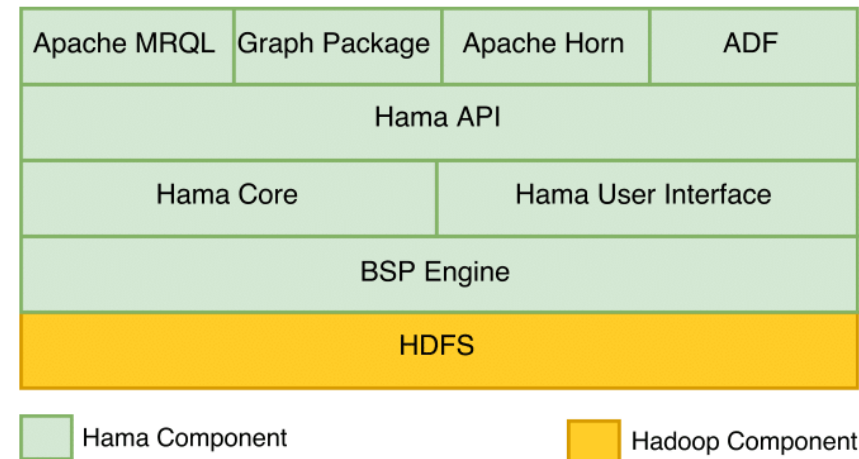
▷ Total runtime of all supersteps

- $W + H.g + S.l$
- where W is the total local comp. time
- g = time to deliver a message
- $H.g$ = total time for communication
- S = total supersteps
- $S.l$ = Total time required for barrier synchronization



Apache Hama

- ▷ Opensource, **in-memory and iterative**
- ▷ Distributed computation framework
- ▷ BSP programming model
- ▷ Processes data stored in HDFS
- ▷ Replacement of MapReduce
- ▷ Can process **massive datasets**
- ▷ **Clean programming interface**



HDFS

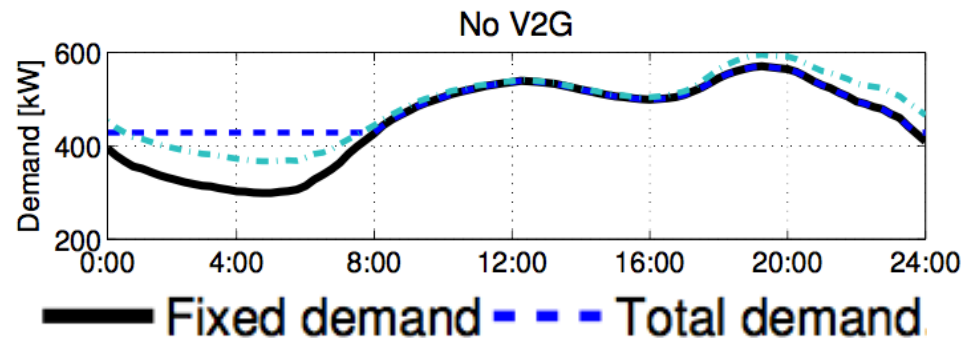
- ▷ *Distributed storage*
- ▷ Runs on *commodity hardware*
- ▷ Master-Slave model (Namenode and DataNode)
- ▷ Files are broken down into blocks of 128MB.
- ▷ Can store huge amounts of data
- ▷ Scalable, available and fault tolerant
- ▷ This thesis : Used to store data

Algorithm

Controlled charging has to strike a balance between

▷ Grid goals (Valley Filling)

- EV charging in valleys of fixed demand
- Avoid peaks
- Leads to stable system

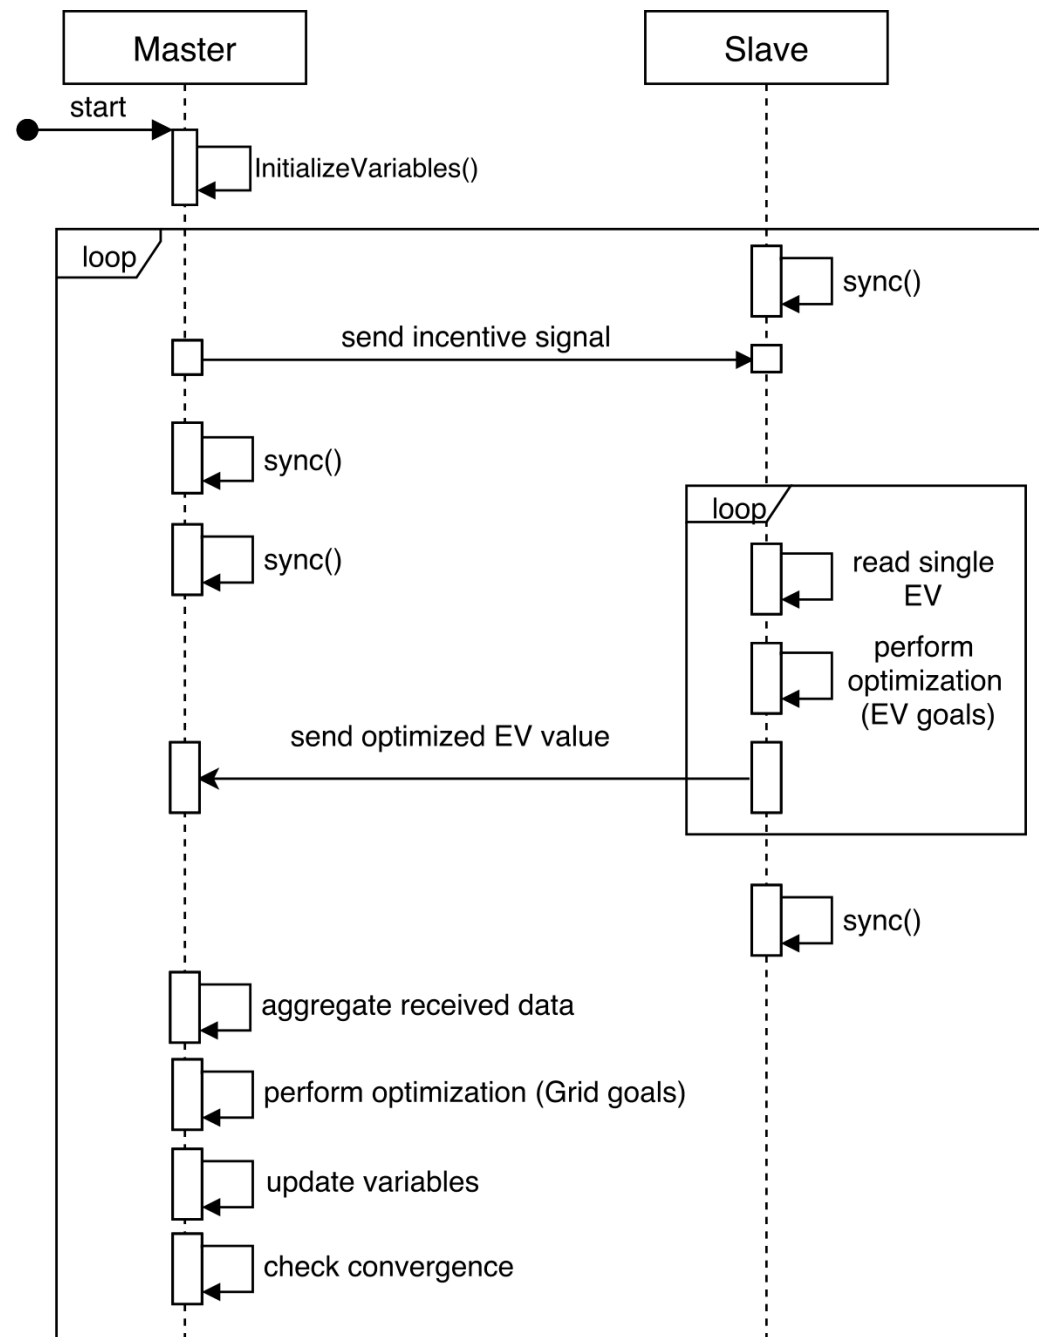


▷ EV goals (EV Optimization problem)

- Minimize battery depreciation

Algorithm

- ▷ Master executes Grid goals
- ▷ Slave executes EV goals
- ▷ sync used for barrier synchronization
- ▷ **1 Master, N Slave tasks**




Deployment



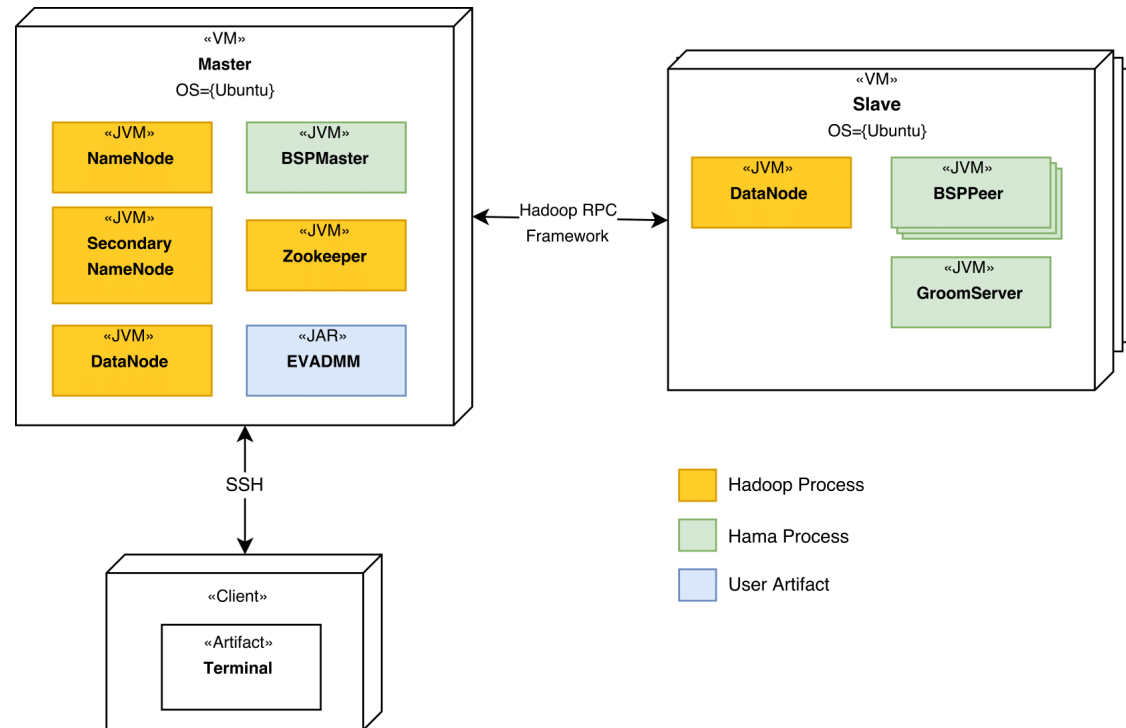
LRZ

- ▷ Leibnitz RechenZentrum
- ▷ Infrastructure as a Service
- ▷ Flexible, secure, highly available
- ▷ Used 10 VMs (Ubuntu, 4 CPU, 8GB RAM)
- ▷ Installed Hadoop, Hama and CPLEX in distributed mode



Deployment

- ▷ Client submits the job to BSPMaster
- ▷ BSPMaster communicates with GroomServer to start multiple BSPPeer (tasks).
- ▷ Each slave can run upto 5-6 tasks in parallel.
- ▷ Example: 9 slaves will have 54 tasks running in parallel executing the submitted job.



Deployment

▷ Client submits the job to

BSPMaster

▷ BSPMaster

GroomServer

BSPPeer (task)

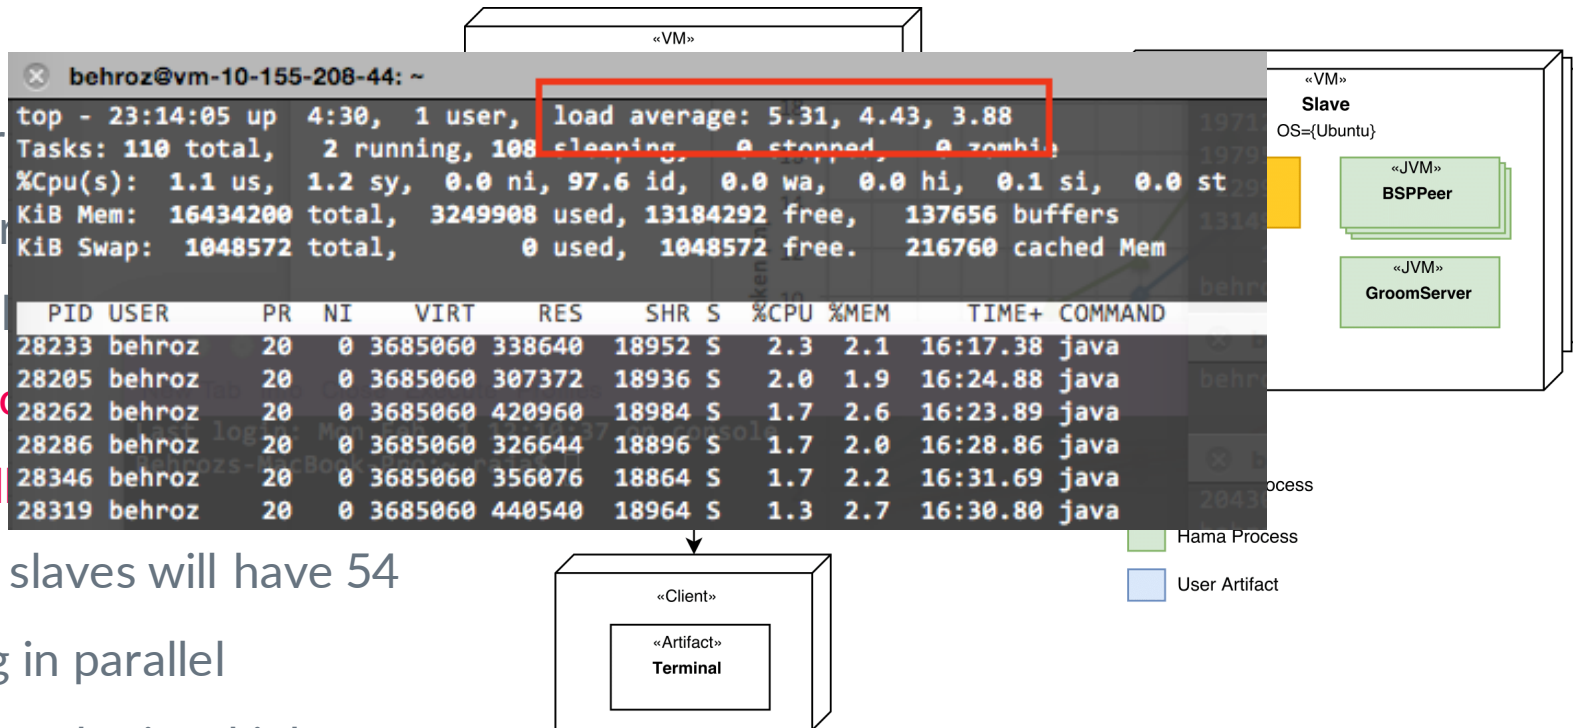
▷ Each slave

tasks in parallel

▷ Example: 9 slaves will have 54

tasks running in parallel

executing the submitted job.



Results


Runtime behavior

- ▷ $T_{total} = W_m + W_s + S.l_s$
- ▷ Time taken while processing 100K EVs for 400 supersteps using 50 tasks.
 - 65% spent on EV optimization (W_s)
 - 25% spent on Master processing (W_m)
 - 10% spent on synchronization ($S.l_s$)

W_s can be decreased by adding more computation power.

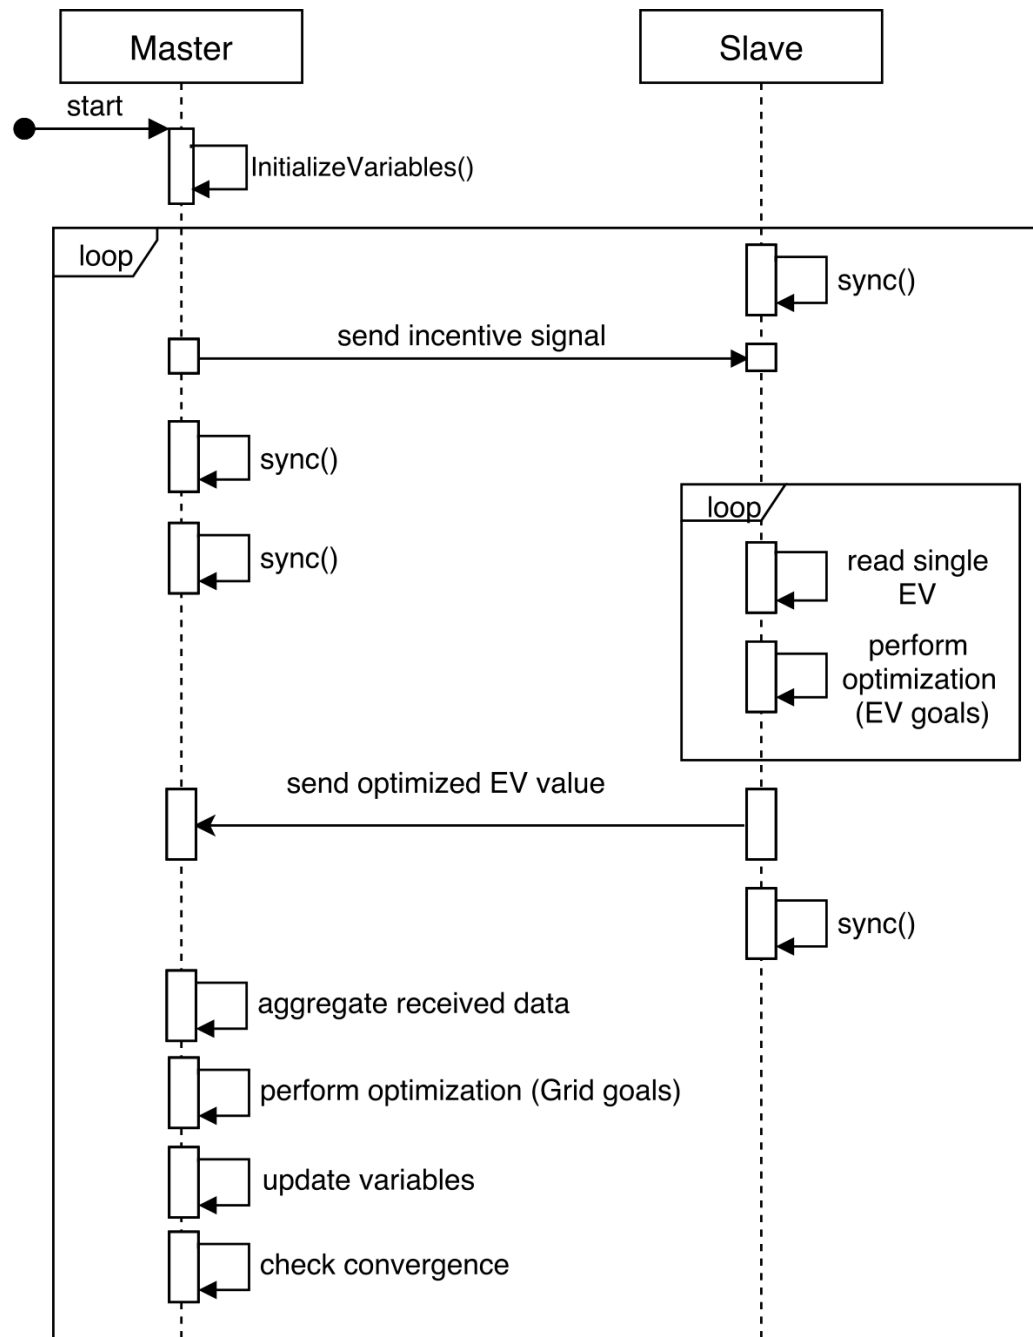
Problem: Processing on Master is a bottleneck.

Solution: Decentralized Approach



Algorithm

- ▷ Measured the time taken by different operations on Master
- ▷ Updated algorithm sends aggregated values only




Decentralized Approach

▷ Total time reduced to 58 minutes from 78.

- 89% spent on EV optimization (W_s)
- **1% spent on Master processing** (W_m)
- 10% spent on synchronization ($S.l_s$)

$$T_{total} = W_s + S.l_s$$



Runtime model

▷ $T_{total} = W_s + S.l_s$

▷ Can be reformulated in terms of total supersteps (S),
total parallel process (P) and total EVs (N)

- t_c represents the time taken to solve a single EV optimization problem. Average value 6.5 ms.
- t_{sync} average time slave spends waiting for barrier synchronization. Average value 57 ms.

$$T_{total} = \frac{S.N}{P}t_c + S.P.t_{sync}$$



Total time to process 1 Mio EVs (N) in 200 supersteps (S) using 100 parallel processes (P) ?

$$T_{total} = \frac{S.N}{P}t_c + S.P.t_{sync}$$

$$T_{total} = 240 \text{ min}^*$$

* Average estimated time



Total machines required to process 1 Mio EVs
(N) in 100 supersteps (S) in 60 minutes (T) ?

$$M = \frac{S.N}{T_{total}.5} t_c$$

$$M = 36^*$$

* Assuming that each machine can run 5 Hama processes in parallel.

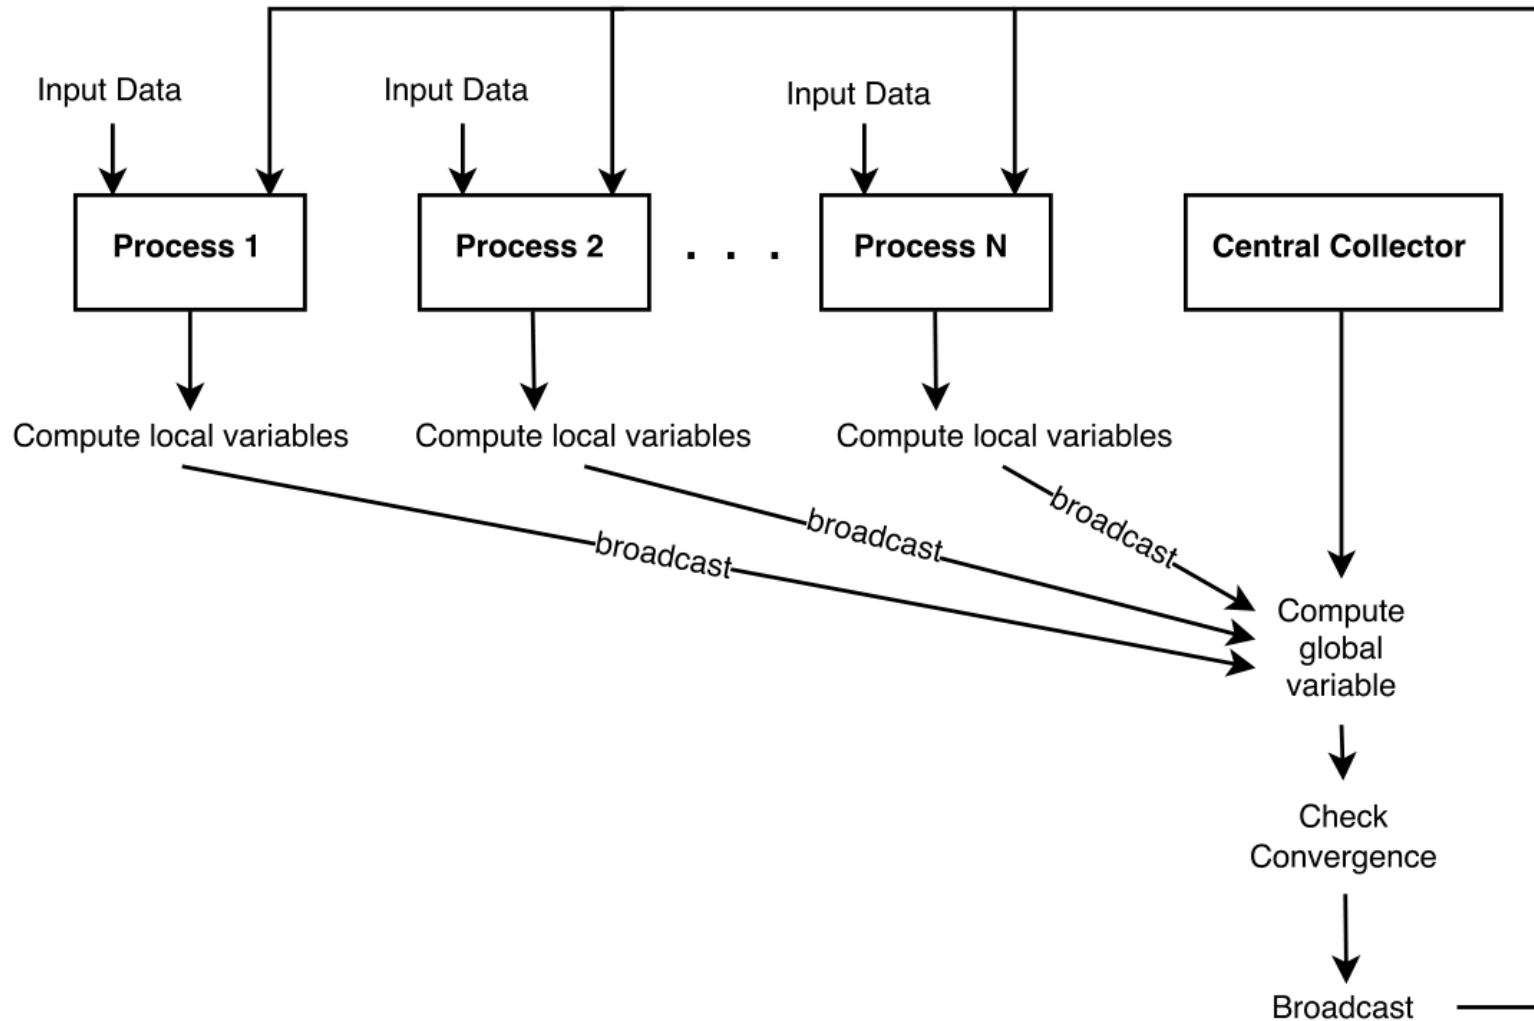
Framework



Framework (ADF)

- ▷ ADMM Distributed Framework
- ▷ For new algorithms, **implementation** needs to be started **from scratch**
- ▷ Reusable set of classes
- ▷ **Automating the optimization process** as much as possible

General Idea



Requirements

- ▷ Solver agnostic
- ▷ Modeling language
- ▷ Supports multiple categories of ADMM problems
- ▷ Functions
- ▷ Abstraction on distributed system

Solver & Modeling Language Agnostic

- ▷ CPLEX, Gurobi, OPL, CPL
- ▷ Exchange problem

$$x_i^{k+1} := \operatorname{argmin}_{x_i} (f_i(x_i) + (\rho/2) \|x_i - x_i^k + \bar{x}^k + u^k\|_2^2)$$
$$u^{k+1} := u^k + \bar{x}^{k+1}$$

- ▷ All we care about is the value of x_i^{k+1}

$$x_i^{k+1} := \text{XUpdate}(\text{function_input}, \text{local_variables})$$

- ▷ User can implement XUpdate interface using any solver or modeling language !

Multiple ADMM Categories

▷ Exchange problem

$$x_i^{k+1} := \operatorname{argmin}_{x_i} (f_i(x_i) + (\rho/2) \|x_i - x_i^k + \bar{x}^k + u^k\|_2^2)$$

$$u^{k+1} := u^k + \bar{x}^{k+1}$$

▷ Consensus problem

$$u_i := u_i + x_i - z$$

$$x_i := \operatorname{argmin}_{x_i} (f_i(x_i) + (\rho/2) \|x_i - z + u_i\|_2^2)$$

$$z := \operatorname{prox}_{g, N_\rho}(\bar{x} + \bar{u})$$

▷ Extend **BSPBase**, **ContextBase** abstract classes

Functions

▷ Exchange problem

$$x_i^{k+1} := \operatorname{argmin}_{x_i} (f_i(x_i) + (\rho/2) \|x_i - x_i^k + \bar{x}^k + u^k\|_2^2)$$
$$u^{k+1} := u^k + \bar{x}^{k+1}$$

▷ **Implement XUpdate** interface for new functions and provide them at startup

Abstraction on distributed system

```
1  ADFJob job = new ADFJob();
2
3  job.setMaxIteration(4);
4  job.setJobName("ADF Exchange EVADMM job");
5  job.setInputPath("aggregator.txt,EVs.txt");
6  job.setOutputPath("output/");
7  job.setSolutionVectorSize(96);
8
9  job.setADMMClass(BSPExchange.class);
10 job.setFunction1(ValleyFillingOptimizationFunction.class);
11 job.setFunction2(EVOptimizationFunction.class);
12
13 job.run();
```

Conclusion

Conclusion

- ▷ **Reproduce EVADMM** on distributed environment
 - Data aggregation on slaves massively decreases the total runtime.
Parallelize time consuming serial parts of algorithm.
- ▷ **Foundation for a general framework to solve similar problems**
- ▷ **Spark vs Hama -> Go with Spark**
 - Performance is equal
 - Hama has very small community
 - Even better -> Use **Apache Beam/Google DataFlow**

Conclusion

- ▷ **Reproduce EVADMM** on distributed environment
 - Data aggregation on slaves massively decreases the total runtime. Parallelize time consuming serial parts of algorithm.
- ▷ Foundation for a general framework to solve similar problems
- ▷ Spark vs Hama -> **Go with Spark**
 - Performance is equal
 - Hama has very small community
 - Even better -> Use **Apache Beam/Google DataFlow**
- ▷ **Contributed to Apache Hama** to add round-robin based task allocation.

Thanks!

Any questions?

You can find me at:
behroz.sikander@tum.de

References

- ▷ S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein. “Distributed optimization and statistical learning via the alternating direction method of multipliers.”
- ▷ S. Boyd and L. Vandenberghe. “Convex programming.”
- ▷ J. Rivera, P. Wolfrum, S. Hirche, C. Goebel, and H.-A. Jacobsen. “Alternating direction method of multipliers for decentralized electric vehicle charging control.”
- ▷ S. Seo, E. J. Yoon, J. Kim, S. Jin, J.-S. Kim, and S. Maeng. “Hama: An efficient matrix computation with the mapreduce framework.”
- ▷ L. G. Valiant. “A bridging model for parallel computation.”

Backup Slides

Hama vs Spark

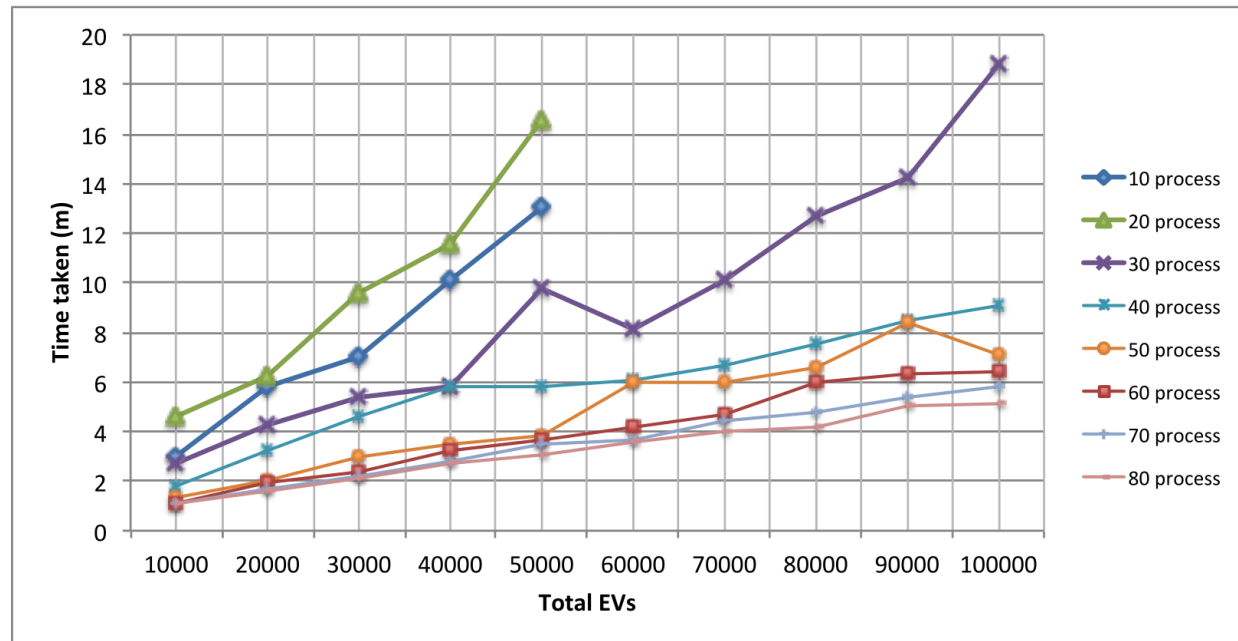
- ▷ From usability perspective, Spark is better than Hama
 - No need to understand the complications of distributed system
- ▷ Hama provides better control on communication and synchronization than Spark
- ▷ Hama is based on strong theoretical backgrounds whereas Spark is still evolving
- ▷ Performance wise, both are more or less equal
 - Recent paper shows that Hama is actually better when performing joins on big datasets
- ▷ Further, ADMM is more natural to implement in BSP.

Lessons

- ▷ Avoid loading data from file system -> In memory
- ▷ Measure load on CPUs early. Overloaded CPU decreases performance
- ▷ Parallel part should be very optimized and carefully implemented. -> Use profiling
- ▷ 3rd party -> Memory leaks

Synchronization Behavior

- ▷ More data -> increase in synchronization time !
- ▷ Reason: Irregular processing time of slaves
- ▷ More processing power -> decreased synchronization time
- ▷ ~10-20% increase in sync. time for adding 10K EVs.



Centralized Master Processing Behavior


▷ Time taken by master to process in 400 supersteps

- x-axis shows total EVs, y-axis shows time taken in minutes
- Each line represents a specific number of parallel tasks used

▷ **More data** means to **more time** process


▷ **Increasing computation power** has **no effect**. Reason: 1 master !

▷ For 100,000 EVs in 400 supersteps it takes **25 minutes !!**



Decentralized Master Processing Behavior

- ▷ Time taken by master to process in 400 supersteps
 - x-axis shows total EVs, y-axis shows time taken in minutes
 - Each line represents a specific number of parallel tasks used
- ▷ No change while increasing data or processing power.
- ▷ Reason: Number of incoming messages to process stays the same !
- ▷ For 100,000 EVs in 400 supersteps it takes 40 seconds!!



Why not 1 Mio ?

- ▷ Unreasonable time taken to process with current cluster

General Idea

▷ Exchange problem

- x-update can be done independently (local variable)
- u-update needs to be computed globally because it depends on all x-updates. (global variable)

$$x_i^{k+1} := \operatorname{argmin}_{x_i} (f_i(x_i) + (\rho/2) \|x_i - x_i^k + \bar{x}^k + u^k\|_2^2)$$
$$u^{k+1} := u^k + \bar{x}^{k+1}$$

▷ Consensus problem

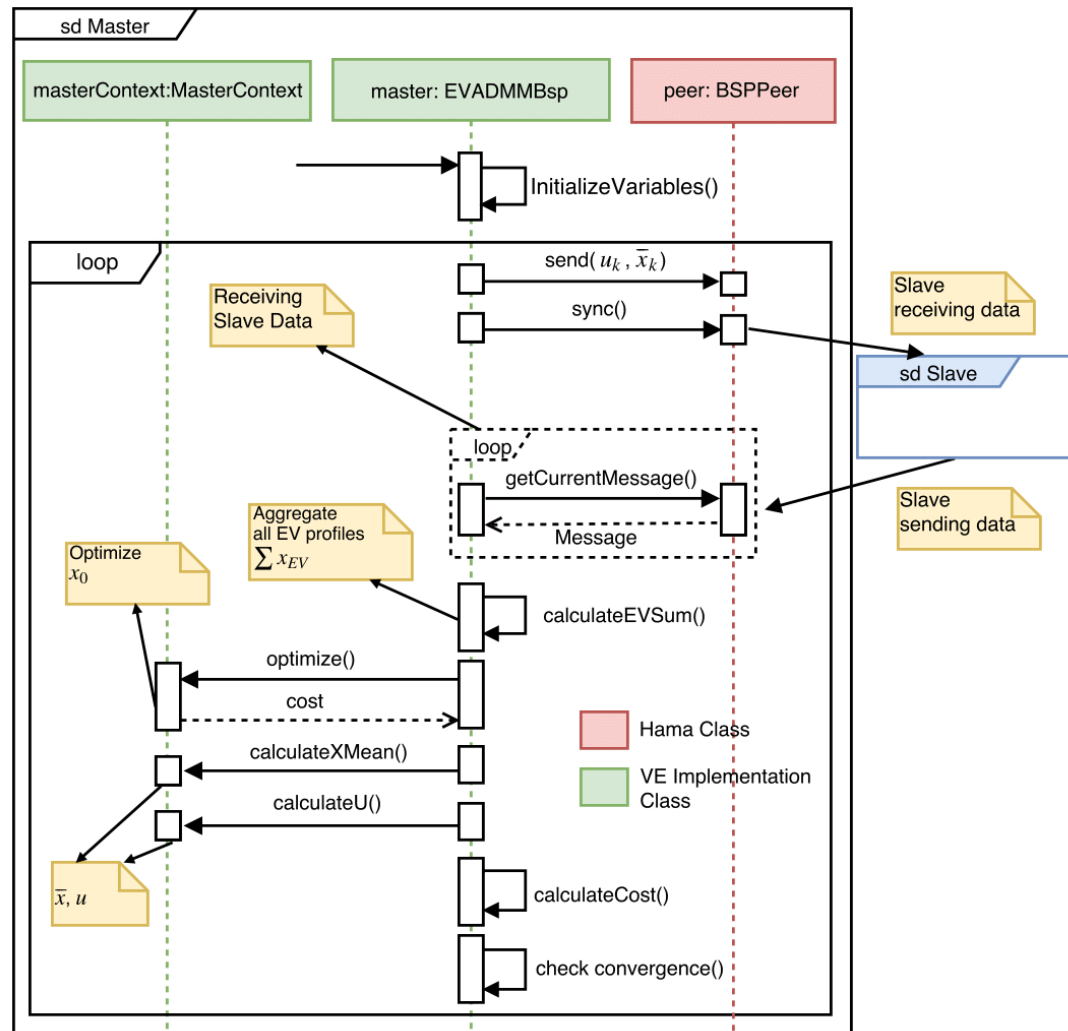
- u and x update can be carried out independently
- z-update can only be computed globally

$$u_i := u_i + x_i - z$$

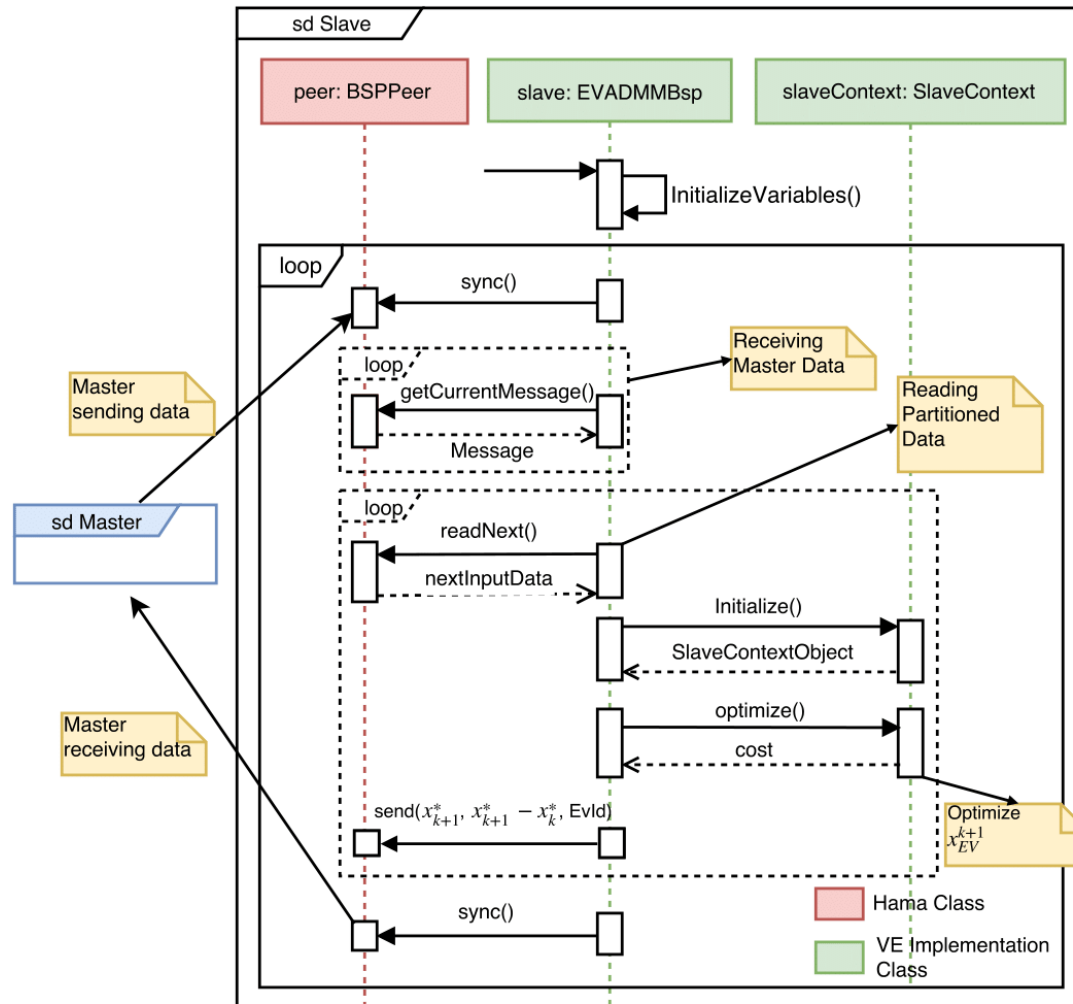
$$x_i := \operatorname{argmin}_{x_i} (f_i(x_i) + (\rho/2) \|x_i - z + u_i\|_2^2)$$

$$z := \operatorname{prox}_{g, N_\rho}(\bar{x} + \bar{u})$$

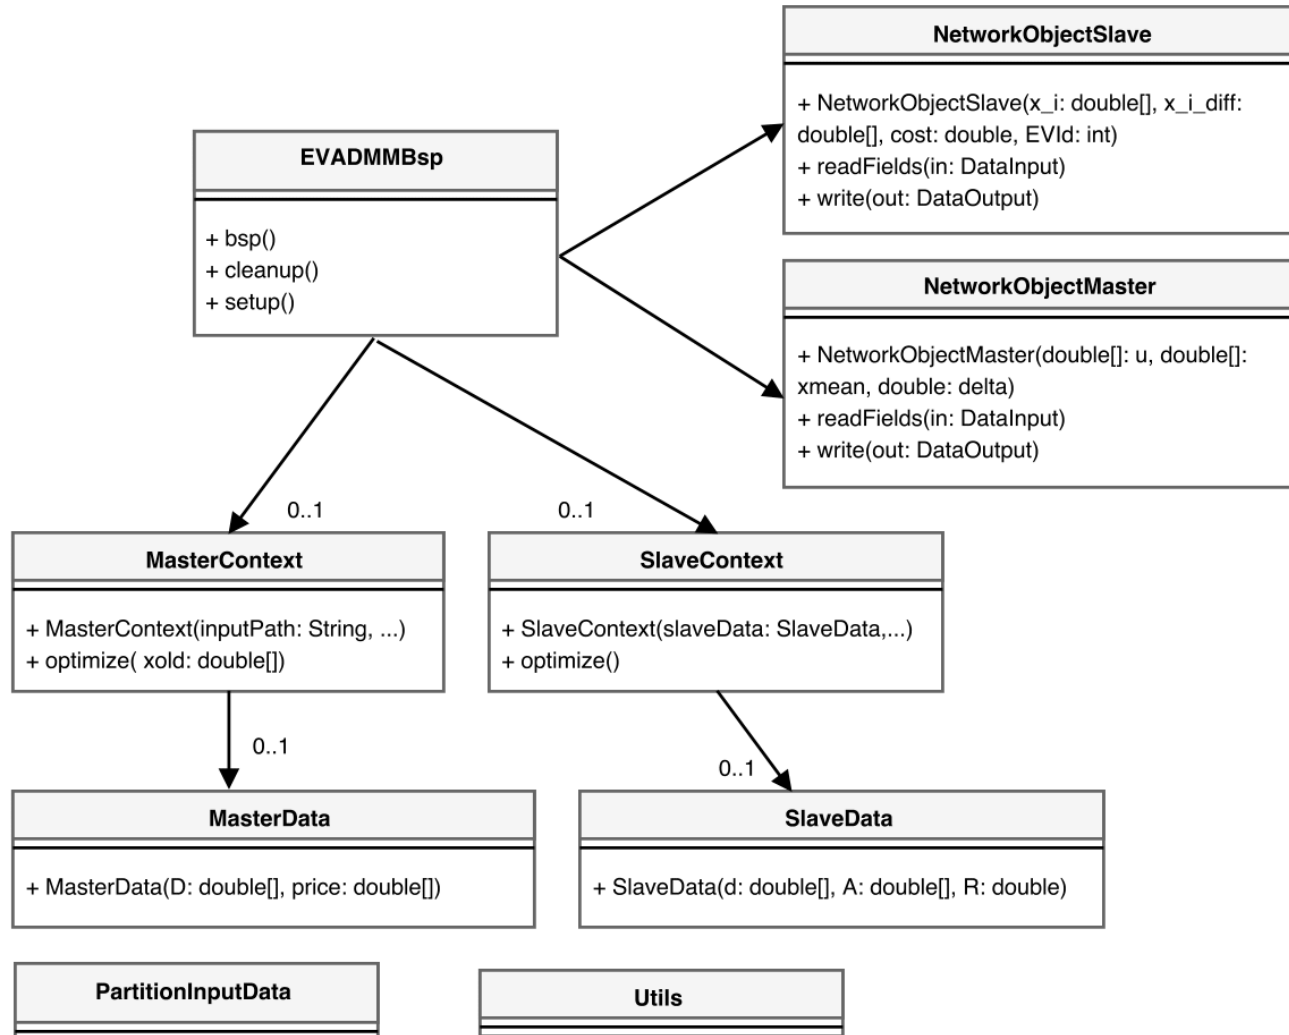
VE- Sequence Diagram



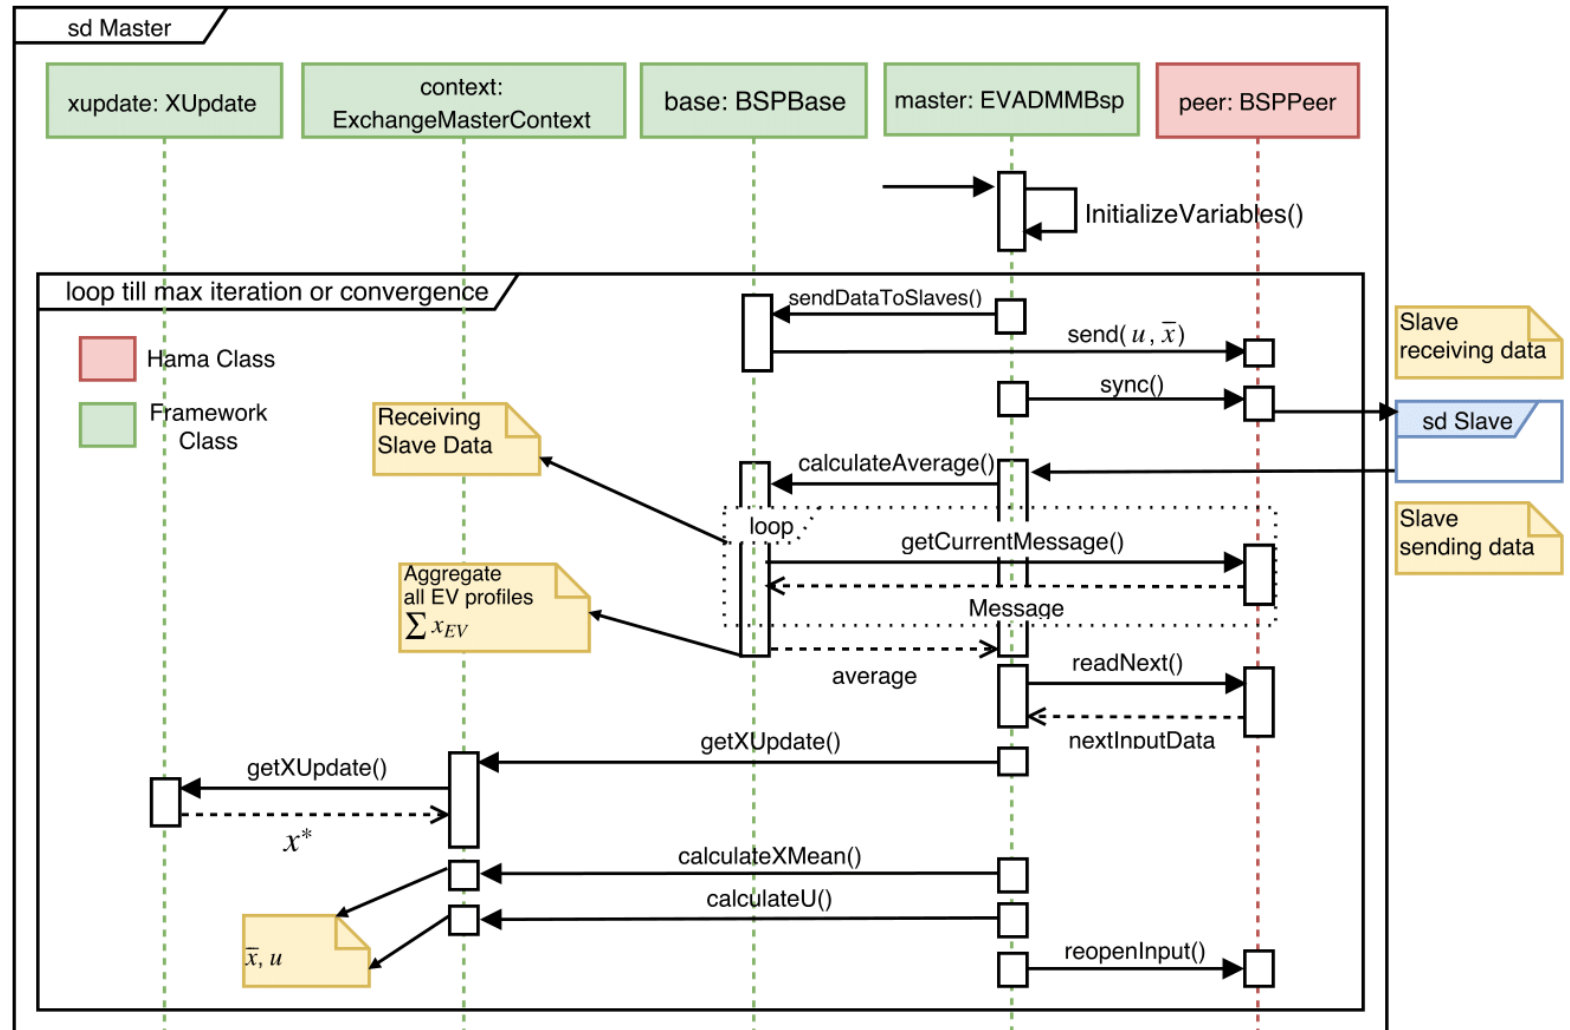
VE- Sequence Diagram



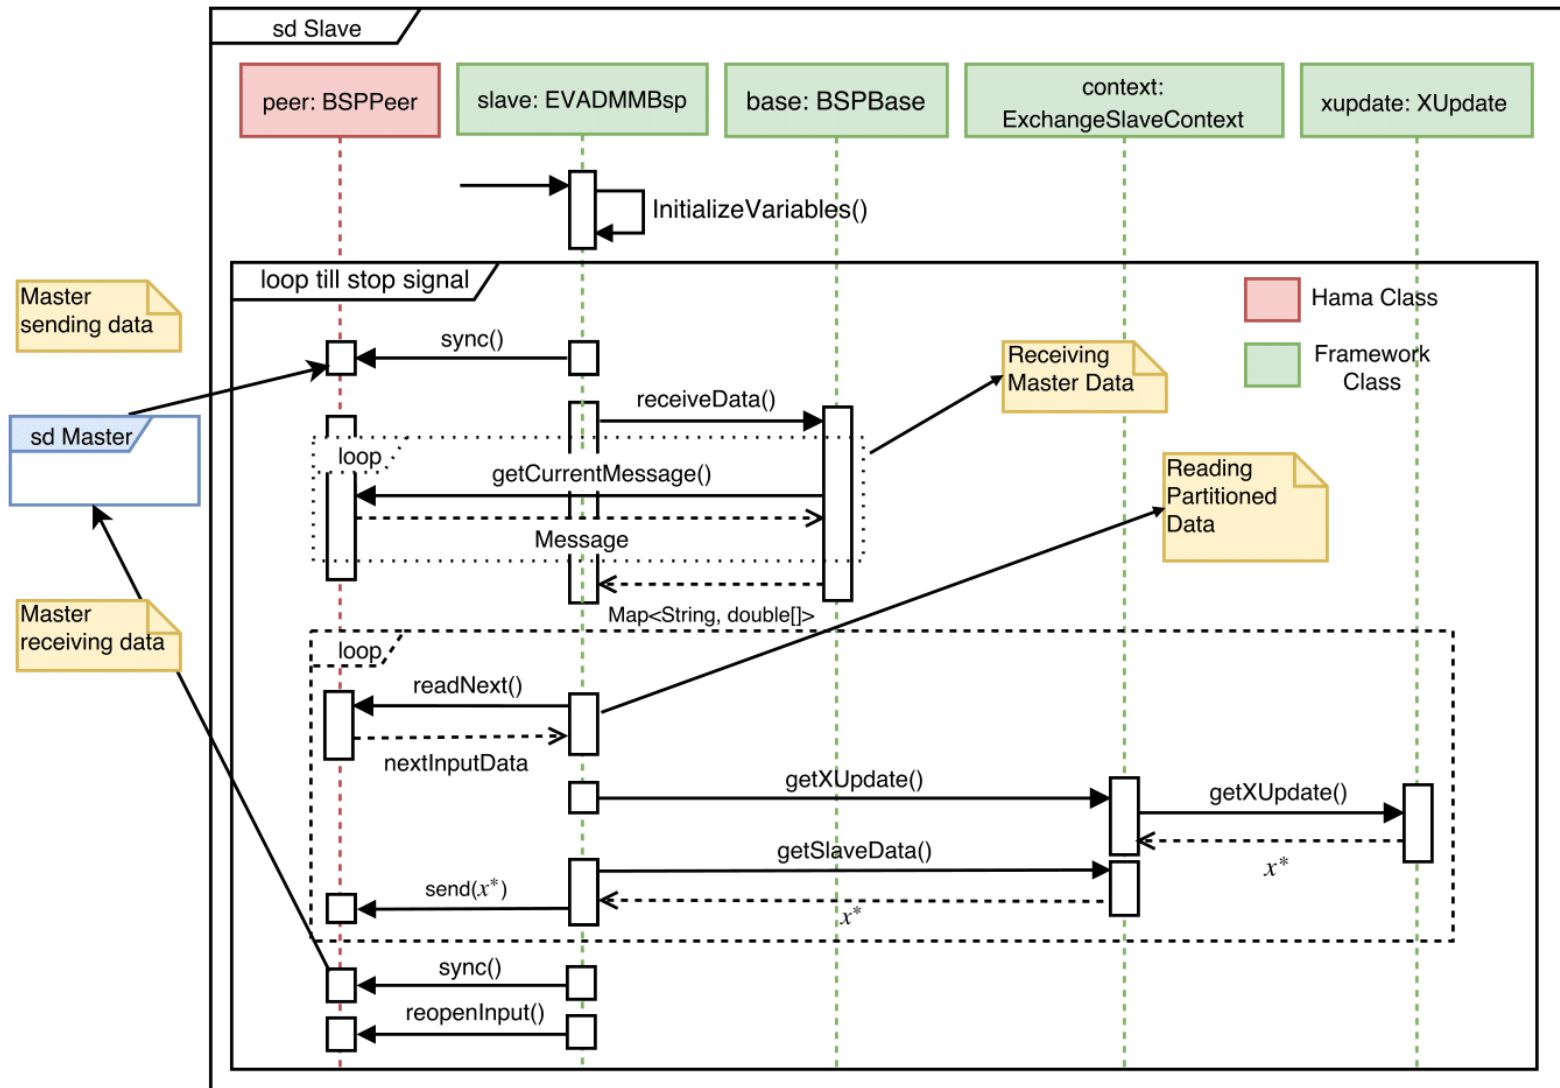
VE- Class Diagram



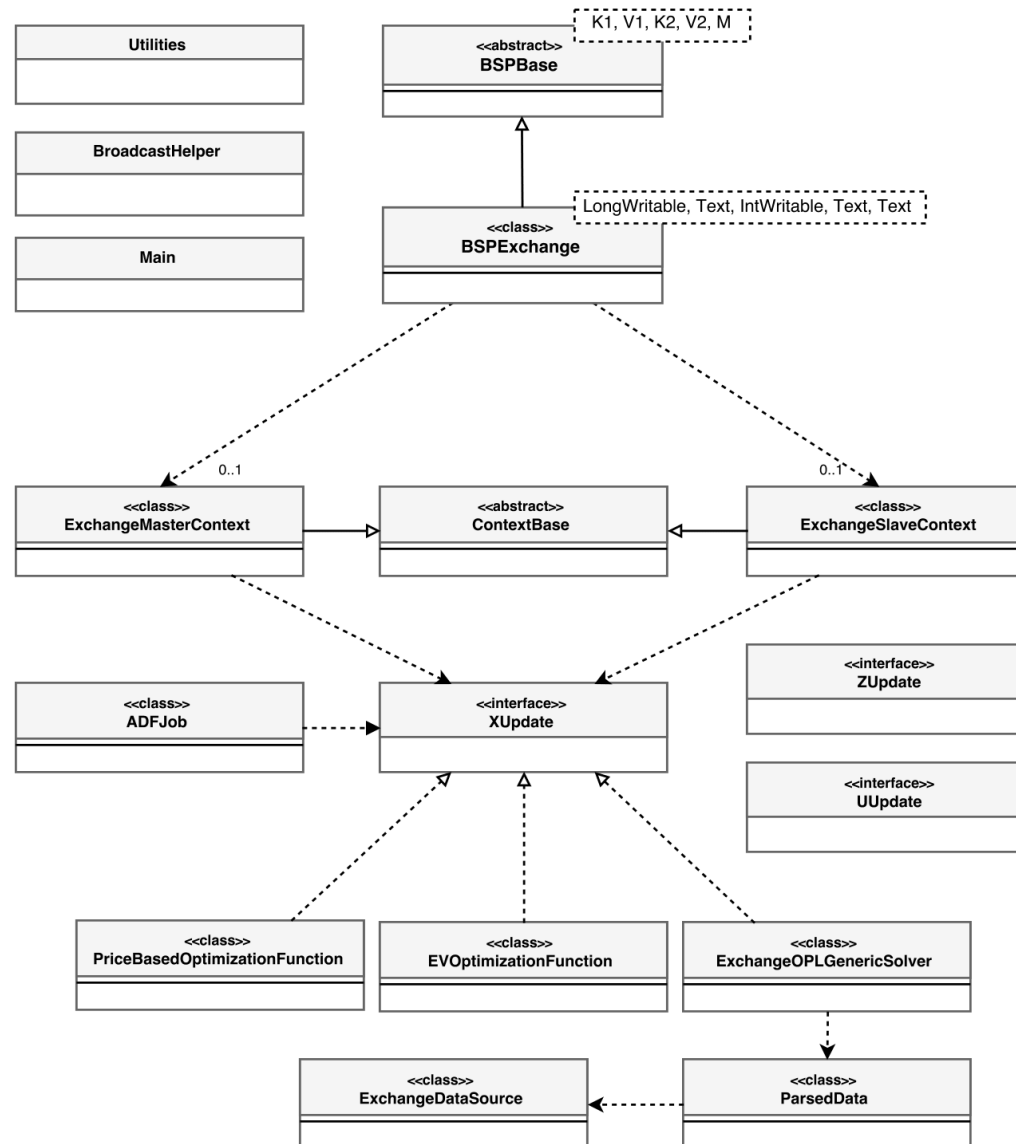
Framework- Sequence Diagram



Framework- Sequence Diagram



Framework- Class Diagram



Background

Optimization

- ▷ **Best solution** from a set of alternatives
- ▷ while being **constrained by a criteria**
- ▷ Examples, portfolio optimization, device sizing

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x), i = 1, \dots, m\end{array}$$

Convex Optimization

- ▷ Objective and constraint functions as **convex** or **concave**
- ▷ ***Solution guaranteed!***

ADMM

- ▷ Alternating Direction Method of Multipliers
- ▷ Used for *distributed* convex optimization
- ▷ Converts problems to local sub-problems
- ▷ Use local solutions to find global solution
- ▷ Iterative

$$\begin{array}{ll}\text{minimize} & f(x) + g(z) \\ \text{subject to} & Ax + Bz = c\end{array}$$

ADMM

▷ Distributed ADMM form

$$x^{k+1} := \operatorname{argmin}_x (f(x) + (\rho/2) \|Ax + Bz^k - c + u^k\|_2^2)$$

$$z^{k+1} := \operatorname{argmin}_z (g(z) + (\rho/2) \|Ax^{k+1} + Bz - c + u^k\|_2^2)$$

$$u^{k+1} := u^k + Ax^{k+1} + Bz^{k+1} - c$$


Solvers

- ▷ Software *tools* to solve mathematical problems.
- ▷ *Abstract out* all the complexities
- ▷ Simple programming interface
- ▷ *Cost functions and constraints can be modeled*
- ▷ Tools: CPLEX, Gurobi, SCIP, CLP, LP_SOLVE
- ▷ This thesis : *CPLEX*

EVADMM

- ▷ Increasing carbon emissions
- ▷ Vehicles are a big source
- ▷ Solution: Electric Vehicles (EV)
- ▷ Charging an EV takes more power than a normal house
- ▷ Solution: Increase infrastructure
- ▷ Or controlled charging

Convergence



Solvers

$$x_i^{k+1} =$$

$$\text{minimize } \rho/2 \|x_i - x_i^k + \bar{x}^k + u^k\|_2^2$$

$$\text{subject to } A_i x_i = R_i$$

$$\underline{x}_i \leq x_i \leq \bar{x}_i$$

```
1 public double optimize(IloCplex cplex){
2     cplex.clearModel();
3     cplex.setOut(null);
4     IloNumVar[] x_i = new IloNumVar[this.x.length];
5     IloNumExpr[] exps = new IloNumExpr[x.length];
6     IloNumExpr[] AXExpEq = new IloNumExpr[x.length];
7
8     double[] data = subtractOldMeanU(this.x);
9
10    for (int i = 0; i < this.x.length; i++) {
11        x_i[i] = cplex.numVar(xi_min[i], xi_max[i]);
12        exps[i] = cplex.prod(rho / 2, cplex.square(cplex.sum(x_i[i],
13        |         |         |         |         |         |         |         |         |
14        |         |         |         |         |         |         |         |         |
15        |         |         |         |         |         |         |         |         |
16        |         |         |         |         |         |         |         |         |
17        |         |         |         |         |         |         |         |         |
18        |         |         |         |         |         |         |         |         |
19        |         |         |         |         |         |         |         |         |
20        |         |         |         |         |         |         |         |         |
21        |         |         |         |         |         |         |         |         |
22        |         |         |         |         |         |         |         |         |
23        |         |         |         |         |         |         |         |         |
24        |         |         |         |         |         |         |         |         |
25        |         |         |         |         |         |         |         |         |
26        |         |         |         |         |         |         |         |         |
27        |         |         |         |         |         |         |         |         |
28        |         |         |         |         |         |         |         |         |
29        |         |         |         |         |         |         |         |         |
30        |         |         |         |         |         |         |         |         |
31        |         |         |         |         |         |         |         |         |
32        |         |         |         |         |         |         |         |         |
33        |         |         |         |         |         |         |         |         |
34        |         |         |         |         |         |         |         |         |
35        |         |         |         |         |         |         |         |         |
36        |         |         |         |         |         |         |         |         |
37        |         |         |         |         |         |         |         |         |
38        |         |         |         |         |         |         |         |         |
39        |         |         |         |         |         |         |         |         |
40        |         |         |         |         |         |         |         |         |
41        |         |         |         |         |         |         |         |         |
42        |         |         |         |         |         |         |         |         |
43        |         |         |         |         |         |         |         |         |
44        |         |         |         |         |         |         |         |         |
45        |         |         |         |         |         |         |         |         |
46        |         |         |         |         |         |         |         |         |
47        |         |         |         |         |         |         |         |         |
48        |         |         |         |         |         |         |         |         |
49        |         |         |         |         |         |         |         |         |
50        |         |         |         |         |         |         |         |         |
51        |         |         |         |         |         |         |         |         |
52        |         |         |         |         |         |         |         |         |
53        |         |         |         |         |         |         |         |         |
54        |         |         |         |         |         |         |         |         |
55        |         |         |         |         |         |         |         |         |
56        |         |         |         |         |         |         |         |         |
57        |         |         |         |         |         |         |         |         |
58        |         |         |         |         |         |         |         |         |
59        |         |         |         |         |         |         |         |         |
60        |         |         |         |         |         |         |         |         |
61        |         |         |         |         |         |         |         |         |
62        |         |         |         |         |         |         |         |         |
63        |         |         |         |         |         |         |         |         |
64        |         |         |         |         |         |         |         |         |
65        |         |         |         |         |         |         |         |         |
66        |         |         |         |         |         |         |         |         |
67        |         |         |         |         |         |         |         |         |
68        |         |         |         |         |         |         |         |         |
69        |         |         |         |         |         |         |         |         |
70        |         |         |         |         |         |         |         |         |
71        |         |         |         |         |         |         |         |         |
72        |         |         |         |         |         |         |         |         |
73        |         |         |         |         |         |         |         |         |
74        |         |         |         |         |         |         |         |         |
75        |         |         |         |         |         |         |         |         |
76        |         |         |         |         |         |         |         |         |
77        |         |         |         |         |         |         |         |         |
78        |         |         |         |         |         |         |         |         |
79        |         |         |         |         |         |         |         |         |
80        |         |         |         |         |         |         |         |         |
81        |         |         |         |         |         |         |         |         |
82        |         |         |         |         |         |         |         |         |
83        |         |         |         |         |         |         |         |         |
84        |         |         |         |         |         |         |         |         |
85        |         |         |         |         |         |         |         |         |
86        |         |         |         |         |         |         |         |         |
87        |         |         |         |         |         |         |         |         |
88        |         |         |         |         |         |         |         |         |
89        |         |         |         |         |         |         |         |         |
90        |         |         |         |         |         |         |         |         |
91        |         |         |         |         |         |         |         |         |
92        |         |         |         |         |         |         |         |         |
93        |         |         |         |         |         |         |         |         |
94        |         |         |         |         |         |         |         |         |
95        |         |         |         |         |         |         |         |         |
96        |         |         |         |         |         |         |         |         |
97        |         |         |         |         |         |         |         |         |
98        |         |         |         |         |         |         |         |         |
99        |         |         |         |         |         |         |         |         |
100    }
101    cplex.addMinimize(cplex.sum(exps));
102    cplex.addEq(cplex.sum(AXExpEq), this.slaveData.getR());
103
104    if (cplex.solve()) {
105        x_optimal = cplex.getValues(x_i);
106    }
107 }
```


Modeling Language

- ▷ Specify equations more *naturally*
- ▷ Internally convert to solver understandable format
- ▷ Model + data file as input

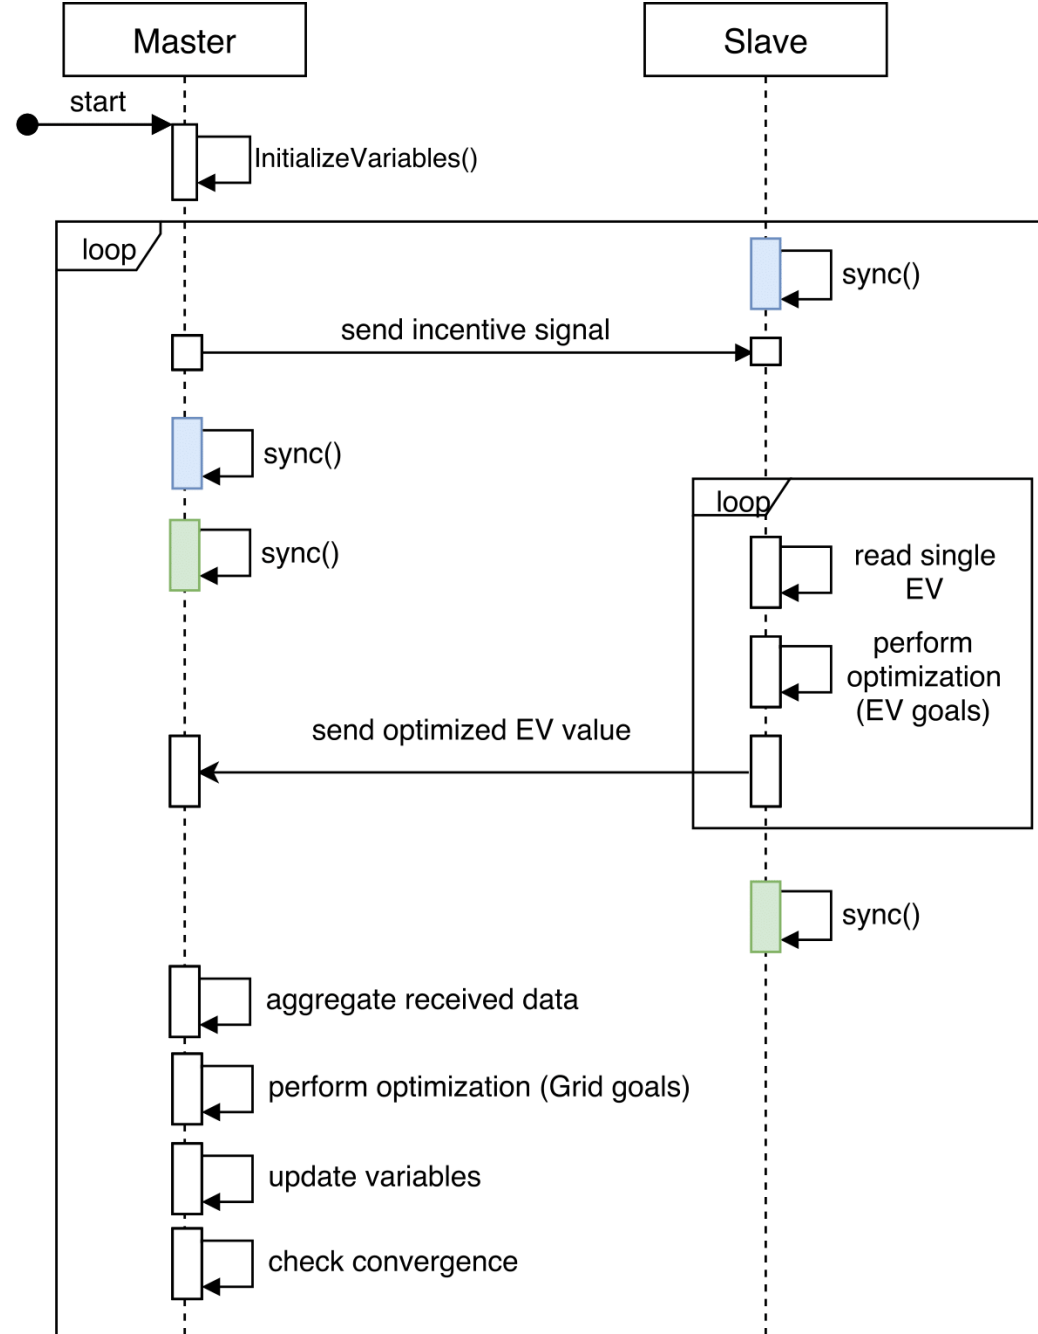
Modeling Language

$$\begin{aligned} x_i^{k+1} = \\ \text{minimize } & \rho/2 \|x_i - x_i^k + \bar{x}^k + u^k\|_2^2 \\ \text{subject to } & A_i x_i = R_i \\ & \underline{x}_i \leq x_i \leq \bar{x}_i \end{aligned}$$

```
1 dvar float xOptimal[i in R] in xi_min[i]..xi_max[i];
2 minimize
3     sum(i in R)
4         (
5             (rho / 2) *
6             (xOptimal[i] - xOld[i] + xMean[i] + u[i]) *
7             (xOptimal[i] - xOld[i] + xMean[i] + u[i])
8         );
9 subject to {
10     sum(i in R)
11         xOptimal[i] * A[i] == R_value;
12 }
```

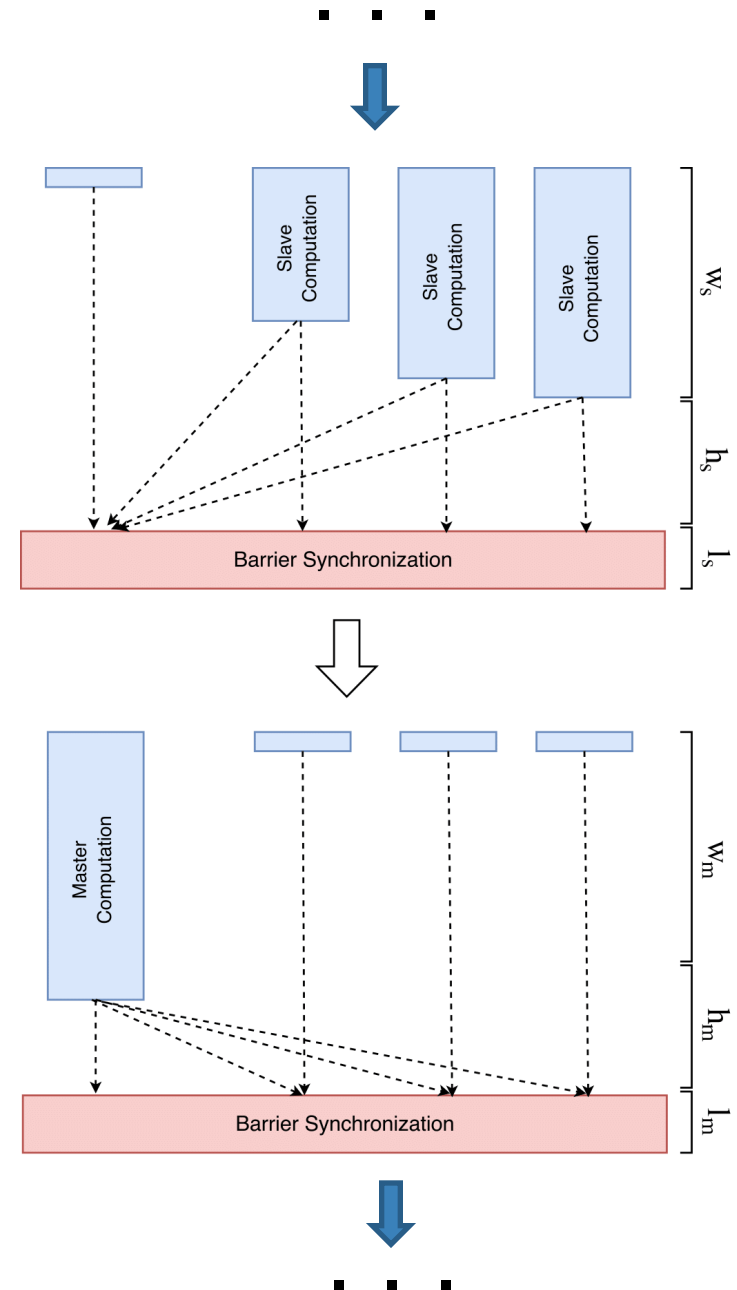
Algorithm

- ▷ 2 supersteps per iteration
- ▷ To process 100K EVs, all Slaves will send 100K messages to Master



Supersteps

- ▷ 2 supersteps per iterations
- ▷ $T_{\text{total}} = T_{\text{slave}} + T_{\text{master}}$
- ▷ Master/slave communication (h_m/h_s) and synchronization time (l_m) are insignificant
- ▷ $T_{\text{total}} = W_m + W_s + S.l_s$



Deployment



Docker

- ▷ light weight, secure, open source project
- ▷ Package application and dependencies in Containers
- ▷ Easily create, deploy and run containers
- ▷ This thesis
 - Deployment on compute server
 - 3 containers
 - HDFS and Hama configured



Docker - Problems

- ▷ /etc/hosts file read only
- ▷ Dynamic Ips
- ▷ No public IP
- ▷ Single harddrive
- ▷ No network latency

Problems- LRZ

- ▷ Memory leaks
- ▷ Processes randomly crashing
- ▷ Firewall issues
- ▷ Hama non-uniform process distribution

Top command

```
behroz@vm-10-155-208-44: ~  
top - 23:14:05 up 4:30, 1 user, load average: 5.31, 4.43, 3.88  
Tasks: 110 total, 2 running, 108 sleeping, 0 stopped, 0 zombie  
%Cpu(s): 1.1 us, 1.2 sy, 0.0 ni, 97.6 id, 0.0 wa, 0.0 hi, 0.1 si, 0.0 st  
KiB Mem: 16434200 total, 3249908 used, 13184292 free, 137656 buffers  
KiB Swap: 1048572 total, 0 used, 1048572 free. 216760 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
28233	behroz	20	0	3685060	338640	18952	S	2.3	2.1	16:17.38	java
28205	behroz	20	0	3685060	307372	18936	S	2.0	1.9	16:24.88	java
28262	behroz	20	0	3685060	420960	18984	S	1.7	2.6	16:23.89	java
28286	behroz	20	0	3685060	326644	18896	S	1.7	2.0	16:28.86	java
28346	behroz	20	0	3685060	356076	18864	S	1.7	2.2	16:31.69	java
28319	behroz	20	0	3685060	440540	18964	S	1.3	2.7	16:30.80	java

Related Work



Serial Implementation

- ▷ Implemented in Matlab
- ▷ Solver: CVXGEN
- ▷ Data processed: 100 EV
- ▷ Total processing time: 2 minutes

Distributed Implementation



- ▷ Implemented in *Matlab*
- ▷ *parfor* function of Parallel Computing Toolbox used
- ▷ 1 machine with 16 cores and 64 GB RAM
- ▷ Data processed: 100,000
- ▷ Solver: CVXGEN
- ▷ Time: *100,000 EVs* processed in *30 minutes*



Overall runtime




Matlab vs Hama Implementation

Change at runtime ?







Algorithm

```
1: procedure VE(max_iter, total_ev,  $\rho$ , input_data_ev, input_data_aggregator,  
2:    $u, \bar{x}, x^* \leftarrow 0$   
3:   for each  $k$  in max_iter do  
4:     for each EV  $i$  in total_ev do  
5:        $data_i \leftarrow$  Read  $i$ -th input from input_data_ev  
6:        $x^* \leftarrow$  PerformEVOptimization( $x^*, u, \bar{x}, data_i$ )  
7:        $X_{EV}[i] \leftarrow$  Store  $x^*$  against EV  $i$   
8:      $sum_{ev} \leftarrow$  Calculate sum of EV profiles  $\sum X_{EV}$   
9:      $data_{agg} \leftarrow$  Read aggregator data from input_data_aggregator  
10:     $x^*_{agg} \leftarrow$  PerformAggregatorOptimization( $x^*_{agg}, u, \bar{x}, data_{agg}$ )  
11:     $\bar{x} \leftarrow (sum_{ev} + x^*_{agg}) / total\_ev$   
12:     $u \leftarrow u + \bar{x}$   
13:    Calculate cost  
14:    if Converged() then  
15:      break
```

Algorithm

```
1: procedure VE(max_iter, total_ev,  $\rho$ , input_data_ev, input_data_aggregator,  
2:    $u, \bar{x}, x^* \leftarrow 0$   
3:   for each  $k$  in max_iter do  
4:     for each EV  $i$  in total_ev do  
5:        $data_i \leftarrow$  Read  $i$ -th input from input_data_ev  
6:        $x^* \leftarrow$  PerformEVOptimization( $x^*, u, \bar{x}, data_i$ )  
7:        $X_{EV}[i] \leftarrow$  Store  $x^*$  against EV  $i$   
8:        $sum_{ev} \leftarrow$  Calculate sum of EV profiles  $\sum X_{EV}$   
9:        $data_{agg} \leftarrow$  Read aggregator data from input_data_aggregator  
10:       $x_{agg}^* \leftarrow$  PerformAggregatorOptimization( $x_{agg}^*, u, \bar{x}, data_{agg}$ )  
11:       $\bar{x} \leftarrow (sum_{ev} + x_{agg}^*) / total\_ev$   
12:       $u \leftarrow u + \bar{x}$   
13:      Calculate cost  
14:      if Converged() then  
15:        break
```

Algorithm

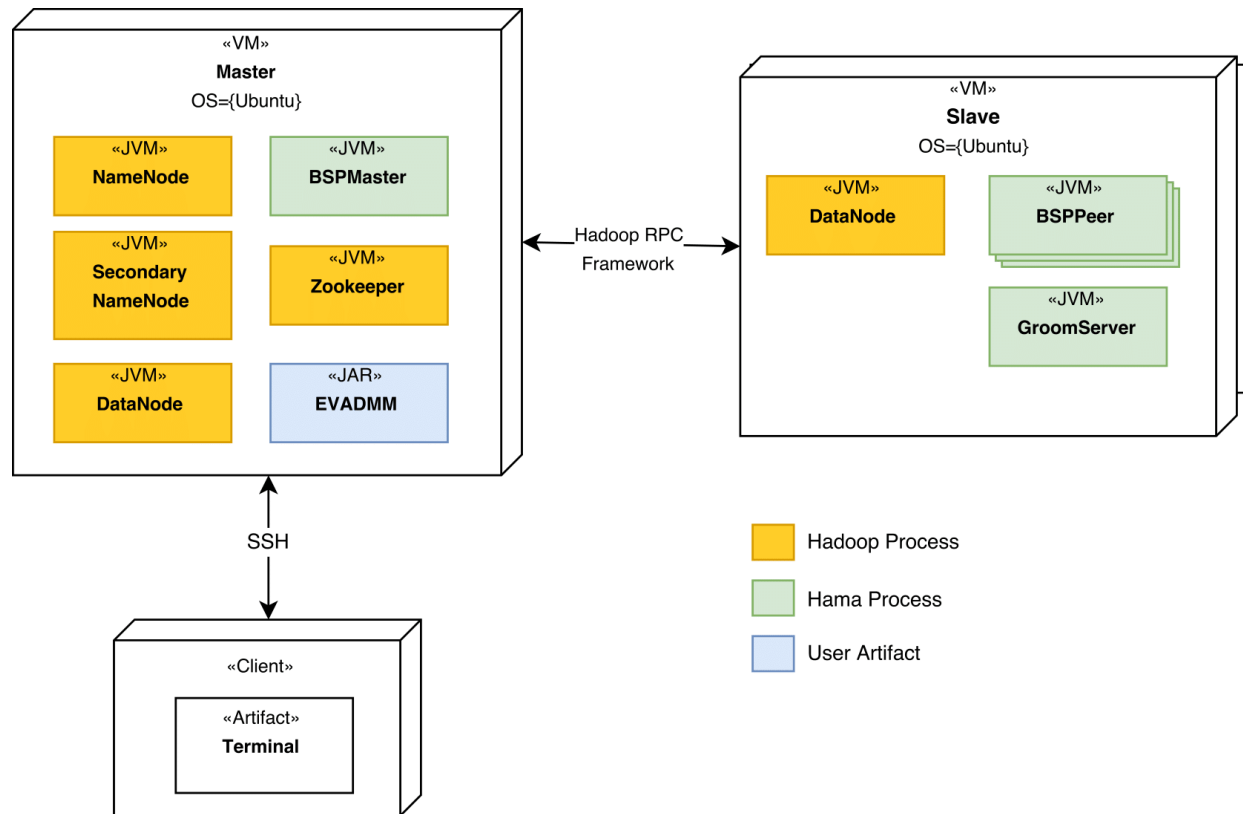
```
1: procedure VE(max_iter, total_ev,  $\rho$ , input_data_ev, input_data_aggregator,  
2:    $u, \bar{x}, x^* \leftarrow 0$   
3:   for each  $k$  in max_iter do  
4:     for each EV  $i$  in total_ev do  
5:        $data_i \leftarrow$  Read  $i$ -th input from input_data_ev  
6:        $x^* \leftarrow$  PerformEVOptimization( $x^*, u, \bar{x}, data_i$ )  
7:        $X_{EV}[i] \leftarrow$  Store  $x^*$  against EV  $i$   
8:      $sum_{ev} \leftarrow$  Calculate sum of EV profiles  $\sum X_{EV}$   
9:      $data_{agg} \leftarrow$  Read aggregator data from input_data_aggregator  
10:     $x^*_{agg} \leftarrow$  PerformAggregatorOptimization( $x^*_{agg}, u, \bar{x}, data_{agg}$ )  
11:     $\bar{x} \leftarrow (sum_{ev} + x^*_{agg}) / total\_ev$   
12:     $u \leftarrow u + \bar{x}$   
13:    Calculate cost  
14:    if Converged() then  
15:      break
```

Runs in **Parallel**
on **Slaves**

LRZ

- ▷ Leibnitz Rechen Zentrum
- ▷ Infrastructure as a Service
- ▷ flexible, secure, highly available
- ▷ All Docker problems gone !

Deployment



Algorithm

```
1: procedure VE(max_iter, total_ev,  $\rho$ , input_data_ev, input_data_aggregator,  
2:    $u, \bar{x}, x^* \leftarrow 0$   
3:   for each  $k$  in max_iter do  
4:     for each EV  $i$  in total_ev do  
5:        $data_i \leftarrow$  Read  $i$ -th input from input_data_ev  
6:        $x^* \leftarrow$  PerformEVOptimization( $x^*, u, \bar{x}, data_i$ )  
7:        $X_{EV}[i] \leftarrow$  Store  $x^*$  against EV  $i$   
8:        $sum_{ev} \leftarrow$  Calculate sum of EV profiles  $\sum X_{EV}$   
9:        $data_{agg} \leftarrow$  Read aggregator data from input_data_aggregator  
10:       $x^*_{agg} \leftarrow$  PerformAggregatorOptimization( $x^*_{agg}, u, \bar{x}, data_{agg}$ )  
11:       $\bar{x} \leftarrow (sum_{ev} + x^*_{agg}) / total\_ev$   
12:       $u \leftarrow u + \bar{x}$   
13:      Calculate cost  
14:      if Converged() then  
15:        break
```

Runs in **Parallel**
on **Slaves**

Algorithm

Abstraction on distributed system

```
1 ADFJob job = new ADFJob();
2
3 job.setMaxIteration(4);
4 job.setJobName("ADF Exchange EVADMM job");
5 job.setInputPath("aggregator.txt,EVs.txt");
6 job.setOutputPath("output/");
7 job.setSolutionVectorSize(96);
8
9 job.setADMMClass(BSPExchange.class);
10 job.setFunction1(ValleyFillingOptimizationFunction.class);
11 job.setFunction2(EVOptimizationFunction.class);
12
13 job.run();
```

Disclaimer: This is an alpha version.